

OpenEC v1.0.0 User Guide

ADSLab @ CUHK

Release: Feb 2019

Contents

1	Preparation	3
2	OpenEC with HDFS-3	5
2.1	Prerequisites	5
2.2	Install HDFS-3 with OpenEC	5
2.3	Example Architecture	6
2.4	HDFS-3 Configuration	6
3	OpenEC with HDFS-RAID	7
3.1	Prerequisites	7
3.2	Install HDFS-RAID with OpenEC	8
3.3	Example Architecture	9
3.4	HDFS-RAID Configuration	9
4	OpenEC with QFS	10
4.1	Prerequisites	10
4.2	Install QFS with OpenEC	11
4.3	Example Architecture	12
4.4	QFS Configuration	12

5	OpenEC Configuration	13
6	Basic Operations	14
6.1	Write	14
6.2	Read	15
6.3	Recovery	16
7	EC Design in OpenEC	16

Abstract

In this user guide, we explain how to install and run OpenEC atop existing distributed storage systems (DSSs). We first explain the preparation steps for running OpenEC (§1). We next explain how to integrate OpenEC with HDFS-3 (§2), HDFS-RAID (§3), and QFS (§4). We then explain how to issue basic operations via OpenEC, including writes, reads (both normal and degraded reads), and recovery. Finally, we show how we can add a new erasure code via OpenEC. Please refer to our FAST’19 paper for the design details of OpenEC.

1 Preparation

OpenEC has been tested in Ubuntu 14.04. We create a user *openec* and install the packages under its home directory `/home/openec`. You may need the `sudo` access in order to install some of the packages via `apt-get`.

Before installing OpenEC, please first install the following prerequisite libraries.

- `cmake` v3.1 or higher

```
$ sudo apt-get install cmake
```

- `g++` v4.8.4

We need a C++ compiler that supports the C++11 standard.

```
$ sudo apt-get install g++
```

- `redis` v3.2.8 or higher

Download and install **redis-3.2.8.tar.gz**.

```
$ tar -zxvf redis-3.2.8.tar.gz
$ cd redis-3.2.8
$ make
$ sudo make install
```

Install `redis` as a background daemon. You can just use the default settings.

```
$ cd utils
$ sudo ./install_server.sh
```

Configure `redis` to be remotely accessible.

```
$ sudo service redis_6379 stop
```

Edit `/etc/redis/6379.conf`. Find the line with `bind 127.0.0.0` and modify it to `bind 0.0.0.0`. Then start redis.

```
$ sudo service redis_6379 start
```

- hiredis

Download and install **hiredis.tar.gz**.

```
$ tar -zxvf hiredis.tar.gz
$ cd hiredis
$ make
$ sudo make install
```

- gf-complete v1.03

Download and install **gf-complete.tar.gz**. Note that you may need to first install `autoconf` and `libtool`.

```
$ tar -zxvf gf-complete.tar.gz
$ cd gf-complete
$ ./autogen.sh
$ ./configure
$ make
$ sudo make install
```

- ISA-L v2.14.0 or higher

Download and install **isa-l-2.14.0.tar.gz**. Note that you may need to first install `yasm`, which is required by ISA-L.

```
$ tar -zxvf isa-l-2.14.0.tar.gz
$ cd isa-l-2.14.0
$ ./autogen.sh
$ ./configure
$ make
$ sudo make install
```

2 OpenEC with HDFS-3

2.1 Prerequisites

The following packages need to be first installed in order to run HDFS-3.

- maven v3.5.0 or higher

Download **apache-maven-3.5.0-bin.tar.gz**.

```
$ tar -zxvf apache-maven-3.5.0-bin.tar.gz
```

Set the environment variables M2_HOME and PATH. You may also need to set MVN_OPTS if you are behind a proxy.

```
export M2_HOME=/home/openec/apache-maven-3.5.0
export PATH=$PATH:$M2_HOME/bin
```

- java8

```
$ sudo add-apt-repository ppa:webupd8team/java
$ sudo apt-get update
$ sudo apt-get install oracle-java8-installer
$ sudo apt-get install oracle-java8-set-default
```

Set the environment variable JAVA_HOME.

2.2 Install HDFS-3 with OpenEC

Download **hadoop-3.0.0-src.tar.gz** (a copy is available on our project website) and extract the source code to /home/openec.

```
$ tar -zxvf hadoop-3.0.0-src.tar.gz
```

We configure the environment variables for HDFS-3. It is recommended to include the following configuration in ~/.bashrc.

```
export HADOOP_SRC_DIR=/home/openec/hadoop-3.0.0-src
export HADOOP_HOME=$HADOOP_SRC_DIR/hadoop-dist/target/hadoop-3.0.0
export PATH=$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
export HADOOP_CLASSPATH=$JAVA_HOME/lib/tools.jar:$HADOOP_CLASSPATH
export CLASSPATH=$JAVA_HOME/lib:$CLASSPATH
export LD_LIBRARY_PATH=$HADOOP_HOME/lib/native:$JAVA_HOME/jre/lib/
amd64/server/:/usr/local/lib:$LD_LIBRARY_PATH
```

Download **openec-v1.0.0.tar.gz** from our project website and extract the source code to `/home/openec`. We can install the patch of OpenEC into HDFS-3 by simply running the script *install.sh*. The script will also compile the modified source code of HDFS-3.

```
$ tar -zxvf openec-v1.0.0.tar.gz
$ cd openec-v1.0.0/hdfs3-integration
$ ./install.sh
```

Please run the following commands to compile OpenEC for HDFS-3.

```
$ cd openec-v1.0.0
$ cmake . -DFS_TYPE:String=HDFS3
$ make
```

2.3 Example Architecture

Table 1 shows an example architecture for our HDFS-3 integration. The OpenEC controller runs in the same node as the HDFS-3 NameNode. Each HDFS-3 DataNode is co-located with an OpenEC agent. Please distribute the working directories (`~/hadoop-3.0.0-src` and `~/openec-v1.0.0`) to all the nodes in the testbed.

IP	HDFS3	OpenEC
192.168.0.1	NameNode	Controller
192.168.0.2	DataNode	Agent
192.168.0.3	DataNode	Agent
192.168.0.4	DataNode	Agent
192.168.0.5	DataNode	Agent

Table 1: Example architecture for HDFS-3 integration.

2.4 HDFS-3 Configuration

We provide sample configuration files under `openec-v1.0.0/hdfs3-integration/conf` for HDFS-3. Here, we show some of the fields related to the integration of OpenEC. You may leave other fields to be the same as our sample configurations. You can copy our sample configuration files to `HADOOP_HOME/etc/hadoop` and configure your HDFS-3 there. Please distribute the configuration files to all the nodes in the testbed.

- `hadoop-env.sh`:

Field	Default	Description
JAVA_HOME	-	Path to java installation. e.g. /usr/lib/jvm/java-8-oracle
HADOOP_CLASSPATH	\$HADOOP_HOME/oeclib/*: \$JAVA_HOME/lib*	Path to OpenEC and java libraries.

- core-site.xml:

Field	Default	Description
fs.defaultFS	hdfs://192.168.0.1:9000	NameNode configuration.
hadoop.tmp.dir	/home/openec/hadoop-3.0.0-src/hadoop-dist/ target/hadoop-3.0.0	Base directory for hdfs3 temporary directories.

- hdfs-site.xml:

Field	Default	Description
dfs.replication	1	Replication factor of HDFS.
dfs.blocksize	1048576	The size of a block in bytes.
dfs.block.replicator.classname	org.apache.hadoop.hdfs.server. blockmanagement. BlockPlacementPolicyOEC	OpenEC placement integration.
link.oec	true	true: Run HDFS3 with OpenEC. false: Run HDFS3 without OpenEC.
oec.controller.addr	192.168.0.1	IP address of OpenEC controller.
oec.local.addr	-	IP address of a node itself.
oec.pktsize	131072	The size of a packet in OpenEC.

- workers:

192.168.0.2
192.168.0.3
192.168.0.4
192.168.0.5

To start HDFS-3, we run the following commands in the NameNode.

```
$ hdfs namenode -format
$ start-dfs.sh
```

3 OpenEC with HDFS-RAID

3.1 Prerequisites

The following packages need first to be installed.

- ant

Download **apache-ant-1.9.13-bin.tar.gz**.

```
$ tar -zxvf apache-ant-1.9.13-bin.tar.gz
```

Set the environment variables ANT_HOME and PATH. You may also need to set ANT_OPTS if you are behind a proxy.

```
export ANT_HOME=~/.apache-ant-1.9.13
export PATH=$PATH:$ANT_HOME/bin
```

- java8

If you have installed java8, please skip this step.

```
$ sudo add-apt-repository ppa:webupd8team/java
$ sudo apt-get update
$ sudo apt-get install oracle-java8-installer
$ sudo apt-get install oracle-java8-set-default
```

Set the environment variable JAVA_HOME.

3.2 Install HDFS-RAID with OpenEC

Download **hadoop-20.tar.gz** (a copy is available on our project website).

```
$ tar -zxvf hadoop-20.tar.gz
```

We configure the environment variables for HDFS-RAID. It is recommended to include the following configuration in ~/.bashrc.

```
export HADOOP_HOME=/home/openec/hadoop-20
export PATH=$HADOOP_HOME/bin:$HADOOP_HOME/sbin:$PATH
export HADOOP_CLASSPATH=$JAVA_HOME/lib/tools.jar:$HADOOP_CLASSPATH
export CLASSPATH=$JAVA_HOME/lib:$CLASSPATH
export LD_LIBRARY_PATH=$HADOOP_HOME/lib/native:$JAVA_HOME/jre/lib/
amd64/server:/usr/local/lib:$LD_LIBRARY_PATH
```

Download **openec-v1.0.0.tar.gz** from our project website and extract the source code to /home/openec. We can install the patch of OpenEC into HDFS-RAID by simply running the script *install.sh*. The script will also compile the modified source code of HDFS-RAID.


```
$ tar -zxvf openec-v1.0.0.tar.gz
$ cd openec-v1.0.0/hdfs3-integration
$ ./install.sh
```

We now compile the source code of OpenEC. Please run the following commands.

```
$ cd openec-v1.0.0
$ cmake . -DFS_TYPE:STRING=HDFSRAID
$ make
```

3.3 Example Architecture

Table 2 shows an example architecture for our HDFS-RAID integration. The OpenEC controller runs in the same node as the HDFS-RAID NameNode. Each HDFS-RAID DataNode is co-located with an OpenEC agent. Please distribute the working directories (`~/hadoop-20` and `~/openec-v1.0.0`) to all the nodes in the testbed.

IP	HDFS-RAID	OpenEC
192.168.0.1	NameNode	Controller
192.168.0.2	DataNode	Agent
192.168.0.3	DataNode	Agent
192.168.0.4	DataNode	Agent
192.168.0.5	DataNode	Agent

Table 2: Example architecture for HDFS-RAID integration.

3.4 HDFS-RAID Configuration

We provide sample configuration files under `openec-v1.0.0/hdfsraid-integration/conf` for HDFS-RAID. Here, we show some of the fields in detail. You may leave other fields to be the same as our sample configurations. You can copy our sample configuration files to `HADOOP_HOME/conf` and configure your HDFS-RAID there. Please distribute the configuration files to all the nodes in the testbed.

- `hadoop-env.sh`:

Field	Default	Description
<code>JAVA_HOME</code>	-	Path to java installation. e.g. <code>/usr/lib/jvm/java-8-oracle</code>
<code>HADOOP_CLASSPATH</code>	<code>\$HADOOP_HOME/oeclib/*:</code> <code>\$JAVA_HOME/lib*</code>	Path to OpenEC and java libraries.

- `core-site.xml`:

Field	Default	Description
fs.default.name	hdfs://192.168.0.1:9000	NameNode configuration.
hadoop.tmp.dir	/home/openec/hadoop-20/tmp	Base directory for HDFS-RAID temporary directories.
topology.script.file.name	-	Path to rackAware.sh

- hdfs-site.xml:

Field	Default	Description
dfs.http.address	192.168.0.1:50070	HTTP address of NameNode
dfs.replication	1	Replication factor of HDFS-RAID.
dfs.block.size	1048576	The size of a block in bytes.
dfs.block.replicator.classname	org.apache.hadoop.hdfs.server.namenode.BlockPlacementPolicyOEC	OpenEC placement class.
link.oec	true	Run HDFS-RAID with OpenEC.
oec.controller.addr	192.168.0.1	IP address of OpenEC controller.
oec.local.addr	-	IP address of a node itself.
oec.pktsize	131072	The size of a packet in OpenEC.

- masters:

192.168.0.1

- slaves:

192.168.0.2
192.168.0.3
192.168.0.4
192.168.0.5

To start HDFS-RAID, we run the following commands in the NameNode.

<pre>\$ hadoop namenode -format \$ start-dfs.sh</pre>

4 OpenEC with QFS

4.1 Prerequisites

The following packages need to be first installed.

- libboost-regex-dev 1.3.4 or higher

```
$ sudo apt-get install libboost-regex-dev
```

- libkrb5-dev

```
$ sudo apt-get install libkrb5-dev
```

- xfslibs-dev

```
$ sudo apt-get install xfslibs-dev
```

- libssl-dev

```
$ sudo apt-get install libssl-dev
```

- python-dev

```
$ sudo apt-get install python-dev
```

- libfuse-dev

```
$ sudo apt-get install libfuse-dev
```

4.2 Install QFS with OpenEC

Download **qfs-v2.1.1.tar.gz** (a copy is available on our project website).

```
$ tar -zxvf qfs-2.1.1.tar.gz
```

We configure the environment variables for QFS. It is recommended to include the following configuration in `~/.bashrc`.

```
export export QFS_HOME=/home/openec/qfs  
export PATH=$QFS_HOME/build/release/bin:$PATH
```

Download **openec-v1.0.0.tar.gz** from our project website and extract the source code to `/home/openec`. We can install the patch of OpenEC into QFS by simply running the script *install.sh*. The script will also compile the modified source code of QFS.

```
$ tar -zxvf openec-v1.0.0.tar.gz
$ cd openec-v1.0.0/qfs-integration
$ ./install.sh
```

We now compile the source code of OpenEC. Please run the following commands.

```
$ cd openec-v1.0.0
$ cmake . -DFS_TYPE:String=QFS
$ make
```

4.3 Example Architecture

Table 3 shows an example architecture for our QFS integration. The OpenEC controller runs in the same node as the QFS metaserver. Each QFS chunkserver is co-located with an OpenEC agent. Please distribute the working directories (`~/qfs` and `~/openec-v1.0.0`) to all the nodes in the testbed.

IP	QFS	OpenEC
192.168.0.1	Metaserver	Controller
192.168.0.2	Chunkserver	Agent
192.168.0.3	Chunkserver	Agent
192.168.0.4	Chunkserver	Agent
192.168.0.5	Chunkserver	Agent

Table 3: Example architecture for QFS integration.

4.4 QFS Configuration

We also provide sample configuration files for QFS. We show some of the fields in detail and other fields can be the same as in our samples.

- MetaServer.conf:

Field	Default	Description
metaServer.clientPort	20000	Port number for metaserver to communicate with clients.
metaServer.chunkServerPort	30000	Port number for metaserver to communicate with chunkserver.
metaServer.logDir	-	Directory for the log of metaserver.
metaServer.cpDir	-	Directory for the checkpoint of metaserver.
openec.useoec	true	true: enable OpenEC integrations. false; disable OpenEC integrations.
openec.localip	-	
openec.coorip	192.168.0.1	IP address of OpenEC controller.

- ChunkServer.conf:

Field	Default	Description
chunkServer.metaServer.hostname	192.168.0.1	IP address of metaserver.
chunkServer.metaServer.port	30000	Port number for metaserver to communicate with chunkserver.
chunkServer.clientPort	22000	Port number for chunkserver to communicate with clients.
chunkServer.chunkDir	-	Directory for chunkserver to store chunks.
chunkServer.pidFile	-	File to store pid of chunkserver
openec.useoec	true	true: enable OpenEC integrations. false; disable OpenEC integrations.
openec.localip	-	
openec.coorip	192.168.0.1	IP address of OpenEC controller.

To start the metaserver, please run the following command:

```
$ metaserver MetaServer.conf
```

To start each chunkserver, please run the following command:

```
$ chunkserver ChunkServer.conf
```

5 OpenEC Configuration

We provide sample configuration files for OpenEC under `openec-v1.0.0/conf`. Table 4 explains the default configuration in our sample. Table 5 and Table 6 show the configuration of an erasure code and an offline encoding pool, respectively.

Field	Default	Description
controller.address	192.168.0.1	IP address of controller.
agents.address	/default/192.168.0.2 /default/192.168.0.3 /default/192.168.0.4 /default/192.168.0.5	A list of IP addresses of all agents, in the form of <i>zone/IP</i> , where <i>zone</i> denotes the zone (e.g. rack or datacenter).
local.address	-	IP address of a node itself.
packet.size	131072	The size of a packet.
dss.type	-	Type of DSS. Please choose from <i>HDFS3</i> , <i>HDFSRAID</i> and <i>QFS</i> .
dss.parameter	-	IP and port of DSS for client access. e.g. <i>192.168.0.1, 9000</i> for HDFS3.
ec.policy		Table 5
offline.pool		Table 6

Table 4: sysSetting.xml for OpenEC

To start OpenEC, we run the following command in controller:

Field	Default	Description
ecid	rs_4_3	Unique id for an erasure code.
class	RSCONV	Class name of erasure code implementation.
n	4	Parameter N for the erasure code.
k	3	Parameter K for the erasure code.
w	1	Parameter W for the erasure code.
opt	-1	Optimization level for OpenEC. Four levels of optimization is provided by OpenEC, including -1, 0, 1, 2. -1: no optimization are enabled. 0: BindX is enabled. 1: BindX and BindY are enabled. 2: Hierarchical awareness is enabled.

Table 5: ec.policy configuration

Field	Default	Description
poolid	rs_4_3_pool	Unique id for an offline encoding pool.
ecid	rs_4_3	Erasure code that is applied for the pool.
base	1	Block size (in MiB) for the pool, which is no larger than the block size in HDFS3.

Table 6: offline.pool configuration

```
$ cd openec-v1.0.0
$ python script/start.py
```

6 Basic Operations

We explain how to issue writes, reads (normal and degraded reads), and recovery via OpenEC.

6.1 Write

OpenEC supports two modes to write a file into a DSS: (i) writing a file with online encoding enabled on the writing path; and (ii) writing a file into an offline encoding pool, in which a coding group is organized and encoded offline.

We run OECClient to issue a write request. We show the usage and a command-line example to write a file (called *input*) of size 3 MiB and store it as */testfile1* with online encoding enabled. The erasure code for online encoding is *rs_4_3* (i.e., RS codes with $n = 4$ and $k = 3$), which is configured in *sysSetting.xml*. Please note that this command should run in a node that holds an Agent.

Usage:

```
./OECClient write [inputfile] [saveas] [ecid] online [sizeinMB]
```

Example:

```
$ ./OECClient write input /testfile1 rs_4_3 online 3
```

We now show how to apply offline encoding. We first write the file into our offline encoding pool. The command-line example in the following shows that we write the file (*input*) and store it as */testfile2*. The offline encoding pool is *rs_4_3_pool*, which is configured in *sysSetting.xml*. Please note that this command should also run in a node that holds an Agent.

Usage:

```
./OECClient write [inputfile] [saveas] [poolid] offline [sizeinMB]
```

Example:

```
$ ./OECClient write input /testfile2 rs_4_3_pool offline 3
```

We then run the following command to instruct OpenEC to start offline encoding.

```
$ ./OECClient startEncode
```

When the offline encoding for a coding group finishes, we can check the log of the controller (*coor_output*). The following line means that the offline encoding for a coding group finishes, where *xxxxxx* denotes the name of the coding group. Note that the name of a coding group is assigned by OpenEC.

```
offlineEnc for xxxxxx finishes
```

6.2 Read

We run OECCClient to read a file, either a normal read or a degraded read. The difference between a normal read and a degraded read is that for a degraded read, some physical blocks in the DSS are unavailable (e.g., deleted) before the read request is issued. Please note that this command should also run in a node that holds an Agent.

Usage:

```
./OECCClient read [filename] [saveas]
```

Example:

```
$ ./OECCClient read /testfile1 output1  
$ ./OECCClient read /testfile2 output2
```

6.3 Recovery

For recovery, we can delete some physical blocks in the DSS and then run the following command to instruct OpenEC to repair the lost blocks.

```
$ ./OECClient startRepair
```

We can see the following information in the log of controller, which denotes that OpenEC finishes repairing a block. Note that xxxxxx is the corresponding block name in OpenEC.

```
repair for xxxxxx finishes
```

7 EC Design in OpenEC

We introduce how to design a new erasure code in OpenEC. The following shows the base class of an erasure code implementation in OpenEC. To add a new erasure code, we need to extend this base class and provide the implementations for *Encode*, *Decode* and *Place* methods. We provide several erasure code implementations under `openec-v1.0.0/src/ec`. Please refer to our sample implementations there.

```
class ECBase {
public:
    int _n, _k, _w;
    int _opt;

    ECBase();
    ECBase(int n, int k, int w, int opt, vector<string> param);

    virtual ECDAG* Encode() = 0;
    virtual ECDAG* Decode(vector<int> from, vector<int> to) = 0;
    virtual void Place(vector<vector<int>>& group) = 0;
};
```